



Machien Learning and Inductive Inference [H02C1a]

Xinhai Zou (r0727971)

Contents

1	Lecture 1: Introduction, Version spaces	1
1.1	Some ML examples in practice	1
1.2	Machine Learning	1
1.3	Machine Learning learning landscape	1
1.4	Some basic concepts and terminology	2
1.5	Input formats (predictive learning)	3
1.6	Output formats, methods (predictive learning)	4
2	Lecture 2: Induction of decision tree	6
2.1	Overview of DT	6
2.2	Learn trees from data	7
2.3	Choosing the best test	7
2.4	Stop splitting nodes	9
2.5	A generic algorithm	10
2.6	Computational complexity	11
2.7	Missing values	13
2.8	Model trees	13
2.9	Multi-target trees	13
2.10	Practical software	14
3	Lecture 3: Learning sets of rules	15
3.1	Linear regression	15
3.2	Trees vs. linear regression vs. inductive bias	15
3.3	Rule learning	16
3.4	Association rules	18
4	Lecture 4: Instance-based learning, Clustering	20
4.1	Instance-based learning	20
4.2	Clustering	22
5	Lecture 5: Evaluating hypotheses	27
5.1	Models and learning algorithms	27
5.2	Statistics	27
5.3	Comparison of evaluation	31
6	Lecture 6: Numerical approaches (ANN, SVM), Computational learning theory	33
6.1	Learning theory	33
7	Lecture 7: Probabilistic approaches, Ensembles	33
8	Lecture 8: Reinforcement learning	33
9	Lecture 9-10: Inductive logic programming	33

1 Lecture 1: Introduction, Version spaces

1.1 Some ML examples in practice

1. Autonomous cars
2. The Robosail project
3. The Robot Scientist
4. Infra Watch, "Hoolandse brug" - the bridge
5. Language learning
6. Automating manual tasks

1.2 Machine Learning

Definition of machine learning: it is the study of how to make programs improve their performance on certain tasks from own (experience).

In this case:

- "performance" = speed, accuracy
- "experience" = earlier observations

Machine Learning vs. other AI

In **machine learning**, the key is **data**, examples of questions and their answer; observations of earlier attempts to solve the problem

In **inductive inference**, it is reasoning from **specific** to **general**, statistics: sample -> population; from **concrete observations** -> **general theory**

1.3 Machine Learning learning landscape

- tasks
 - clustering
 - classification
 - regression
 - reinforcement learning
- techniques
 - Convex optimization
 - Matrix factorization
 - Transfer learning
 - Learning theory
 - Greedy search

- models
 - automata
 - neural network
 - deep learning
 - statistical relational learning
 - decision trees
 - support vector machines
 - nearest neighbors
 - rule learners
 - bayesian learning
 - probabilistic graphical models
- applications
 - natural language processing
 - vision
 - speech
- related courses
 - neural computing
 - support vector machine
 - uncertainty in AI
 - data mining
 - genetic algorithms and evolutionary computing

1.4 Some basic concepts and terminology

- Predictive learning
 - Definition: learn a model that can predict a particular property/ attribute/ variable from inputs
 - Binary classification: distinguish instances of class C from other instances
 - Classification: assign a class C (from a given set of classes) to an instances
 - Regression: assign a numerical value to an instance
 - multi-label classification: assign a set of labels (from a given set) to an instance
 - multivariate regression: assign a vector of numbers to an instances

- multi-target prediction: assign a vector of values (numerical, categorical) to an instances
- Descriptive learning
 - Definition: given a dataset, describe certain patterns in the dataset, or in the population it is drawn from
- Typical tasks in ML
 - function learning: learn a function $X \rightarrow Y$ that fits the given data
 - distribution learning: distribution learning
 - * parametric: the function family of the distribution is known, we only need to estimate its parameters
 - * non-parametric: no specific function family assumed
 - * generative: generate new instances by random sampling from it
 - * discriminative: conditional probability distribution
- Explainable AI (XAI)
 - Definition: means that the decisions of an AI system can be explained
 - Two different levels:
 - * We understand the (learned) model
 - * We understand the individual decision

1.5 Input formats (predictive learning)

- Set
 - training set: a set of examples, instance descriptions that include the target property (a.k.a. labeled instances)
 - prediction set: a set of instance descriptions that do not include the target property ('unlabeled' instances)
 - prediction task: predict the label of the unlabeled instances
- Outcome of learning process
 - transductive learning: the predictions themselves
 - inductive learning: a function that can predict the label of any unlabeled instance
- Explainable AI
 - interpretable: can be interpreted
 - black-box: non-interpretable
- Learning

- Supervised learning: from labeled
- Unsupervised learning: from unlabeled
- Semi-supervised learning: from a few labeled and many unlabeled
- Format of input data
 - input is often assumed to be a set of instances that are all described using the same variables (features, attributes)
 - i.i.d.: independent and identically distributed
 - * tabular data (NN)
 - * sequences
 - * trees
 - * graph
 - * raw data: learning meaningful features from raw data
 - * knowledge: inductive logic programming

1.6 Output formats, methods (predictive learning)

The **output** of a learning system is a model.

- output
 - parametrized functions
 - conjunctive concepts: a conjunctive concept is expressed as a set of conditions, all of which must be true
 - rule sets (if...then...else...)
 - decision trees
 - neural networks
 - probabilistic graphical models
- search methods
 - discrete spaces - methods: hill-climbing, best-first
 - continuous spaces - methods: gradient descent
- typically
 - model structure not fixed in advance - discrete
 - fixed model structure, tune numerical parameters - continuous
- hypothesis space
 - definition: all possible instances
 - for robot example: $\{B,R,M,?\} \times \{S,T,?\} \times \{L,W,?\} \times \{1,2,?\}$

- Version space
 - using candidate elimination
 - pros
 - * can be used for discrete hypothesis spaces
 - * search for all solutions, rather than just one, in an efficient manner
 - * importance of generality ordering
 - cons
 - * not robust to noise
 - * only conjunctive concepts

2 Lecture 2: Induction of decision tree

2.1 Overview of DT

- A decision tree represents a decision procedure where
 - you start with one question
 - the answer will determine the next question
 - and repeat, until you reach a decision
- We will usually call the questions "tests" and the decision a "prediction"
- attribute
 - input attribute $X = \{X_1, X_2 \dots, X_n\}$
 - target attribute Y
 - the tree represents a function $f: X \rightarrow Y$
- Example: Playing Tennis Tree
 - Outlook: $X_1 = \{\text{Sunny, Overcast, Rainy}\}$
 - Humidity: $X_2 = \{\text{High, Normal}\}$
 - Wind: $X_3 = \{\text{Strong, Weak}\}$
 - Tennis: $Y = \{\text{Yes, No}\}$
 - The tree represents a function $\text{Outlook} \times \text{Humidity} \times \text{Wind} \rightarrow \text{Tennis}$
- Boolean tree
- Continuous input attributes
 - We cannot make a different child node for each possible value!
 - Solution: use comparative test \rightarrow a finite number of possible outcomes
- Type of trees
 - target attribute Y is nominal \rightarrow classification tree
 - target attribute Y is numerical \rightarrow regression tree
- **Advantages of Tree (Why tree?)**
 - Learning and using tree is **efficient**
 - Tend to have **good predictive accuracy**
 - Tree is **interpretable**

2.2 Learn trees from data

- Two tasks for DT
 - Task 1: find the smallest tree T such that $\forall (x, f(x)) \in D: T(x) = f(x)$ (meaning that only fulfill current data set)
 - Task 2: find the tree T such that for x drawn from population D , $T(x)$ is (on average) maximally similar to $f(x)$ (T : model tree from data set D , $f(x)$: true function in population D)
 - * loss function: $l: Y_1 \times Y_2 \rightarrow R$ (where Y_1 is predicted value, Y_2 is actual value)
 - * risk R of T , the expectation of loss function, is $E_{x \sim D}[l(T(x), f(x))]$, which is needed to be minimal.
- the basic principle
 - The approach is known as "Top-down induction of decision trees (TDIDT)", or "recursive partitioning"
 - * 1. start with the full data set D
 - * 2. find a test such that examples in D with the same outcome for the test tend to have the same value of Y
 - * 3. split D into subsets, one for each outcome of that test
 - * 4. repeat this procedure on each subset that is not yet sufficiently "pure" (meaning, not all elements have the same Y)
 - * 5. keep repeating until no further splits possible
- rule representation of tree
 - trees can be written as if-then-else rules
 - rules can be simplified
- Two main questions?
 - How to choose which test should be the first? (guess: attribute with minimal entropy?)
 - When to stop splitting nodes? (guess: till pure? or threshold for probability?)

2.3 Choosing the best test

- We focus on classification tree (Y is nominal)
- Information theory
 - a good test is a test that carries much information about the class
 - "entropy" or "missing information": how many bits needed, on average, to convey a piece of information, if an optimal encoding is used

- **bits** <- **Question!! Do not understand the bit, how is it related to entropy?**
- But whatever, the cleverest bit has been provided as below, which is called "entropy"
 - * $e = - \sum_{i=1}^k p_i \log_2 p_i$
- The number e reflects the minimal number of bits that you will need, on average, to encode one value. It is the inherent information content, or **entropy**.
- for classification - we use class entropy
 - The class entropy of a set S of objects (x,y) , where y can be any of k classes c_i , is defined as
 - * $CE(S) = - \sum_{i=1}^k p_i \log_2 p_i$ with $p_i = \frac{(|\{(x,y) \in S | y=c_i\}|)}{|S|}$
 - * it measures how much uncertainty there is about the class of a particular instance
 - high entropy = "many possibilities, all equally likely" - not good for splitting
 - low entropy = "few possibilities, safer to conclude" - better for splitting nodes
 - entropy measures uncertainty
- information gain (IG)
 - will have the same effect of attribute entropy
 - the information gain of a question = the expected reduction of entropy by obtaining the answer to the question
 - in the case of classification trees: expected reduction of class entropy:
 - * $IG(S, t) = CE(S) - \mathbb{E}(CE(S_i)) = CE(S) - \sum_{i=1}^o \frac{|S_i|}{|S|} CE(S_i)$
 - * with t a test, o the number of possible outcomes of attribute/test t , and S_i the subset of S for which the i 'th outcome was obtained.

However, if we focus on regression tree (Y is numerical), can we still use class entropy or information gain?

- Now we assume Y is numerical
- Now we use **variance reduction** instead of entropy or information gain
 - the variance of Y in a set S of instances (x,y) is
 - $Var(S) = \frac{\sum_{(x,y) \in S} (y - \hat{y})^2}{|S| - 1}$ with $\hat{y} = \frac{\sum_{(x,y) \in S} y}{|S|}$
 - and the variance reduction will be (which is similar to information gain (IG))
 - $VR(S, t) = Var(S) - \sum_{i=1}^o \frac{|S_i|}{|S|} Var(S_i)$

2.4 Stop splitting nodes

- In principle, keep splitting until all instances in a subset have the same Y value
 - However, this is useful for task 1, but less useful for task 2 (may cause overfitting problem!)
 - Please remember this is not population, this is just a sampling sample.
- overfitting
 - overfitting improves the consistency of the tree with the given data set D , but may decrease accuracy for instances outside D (whole population \mathcal{D})
- How to avoid overfitting?
 - "cautious splitting": do not split a node unless you are certain that the split is meaningful
 - "post-pruning": do not bother about overfitting while splitting nodes, but once the (large) tree has been built, prune away branches that turned out not to contribute much
- "Cautious splitting"
 - How do we know when not to split a node any further?
 - * a simple approach: try to guess when accuracy on unseen data is going down ("**the turning point**")
 - But how can we guess this turning point, if the data is unseen?
 - * Solution: we can use "validation data" to evaluate
 - * this "validation data" is not for growing tree, only for estimating accuracy on "unseen" data
- "Post-pruning"
 - What if the accuracy will increase again? - afraid to end the growth too early
 - * Solution: we can compute the whole data set, then find its highest point
 - Principle
 - * grow the tree to its full size, then cut away branches that did not contribute to getting better predictions
 - * How to decide which branch does not contribute
 - check for each node in the tree, starting at the bottom: "what would be the accuracy if I cut the tree here?" - if that accuracy is not lower, cut the tree, otherwise, do not cut.

- Pros and Cons
 - post-pruning requires more effort to build a large tree, while it gives more accurate tree
 - cautious splitting is more efficient, while the accuracy is not as good as post-pruning

2.5 A generic algorithm

- TDIDT = "Top-down induction of decision trees", also referred to as "recursive partitioning"
- most decision tree learners follow the same basic approach, but differ in the details
 - we will have a look at the commonalities and differences

```

function TDIDT(E: set of examples) returns tree;
  T' = grow_tree(E);
  T = prune_tree(T');
  return T;

function grow_tree(E: set of examples) returns tree;
  T = generate_tests(E);
  t = best_test(T, E);
  P = {E1, E2, ..., Ek} with Ei = {x ∈ E | t(x) = vi}
  P = partition induced on E by t
  if stop_criterion(E, P)
  then return leaf(info(E))
  else
    for all Ei in P: Ti := grow_tree(Ei);
    return node(t, {(v1, T1), (v2, T2), ..., (vk, Tk)});

```

Figure 1: Algorithm for "Top-down induction of decision trees"

- The blue functions are where implementations differ
 - `prune_tree`: how to prune the tree afterward
 - `generate_tests`: which tests to consider
 - `best_test`: which test to select (use heuristics: entropy or variance)
 - `stop_criterion`: when to stop (cautious splitting or post-pruning)
 - `info`: what information to store in the leaf
- `generate_tests`
 - for numerical attributes: oblique trees can be used for determining c(threshold), while it will be more difficult even though it has higher accuracy. Thus, in practice, non-oblique trees are much more common.

- [best_test](#)
 - for classification trees: information gain (IG) = reduction to entropy
 $CE(S) = -\sum_{i=1}^k p_i \log_2 p_i$
 - * in some cases: "Gini impurity" instead of entropy: $Gini(S) = 1 - \sum_{i=1}^k p_i^2$
 - for regression trees: reduction of variance σ , or reduction of standard deviation $\sqrt{\sigma}$
 - * in some cases: normalization is implemented by "Split information" (SI): $SI(S, t) = -\sum_{i=1}^n \frac{|S_i|}{|S|} \log_2 \frac{|S_i|}{|S|}$
 - * and the Gain Ratio (GR) will be: $GR(S, t) = \frac{IG(S, t)}{SI(S, t)}$
 - for numerical inputs, how to determine c?
 - * Solution: typically, all values are tried, and the c that yields the best heuristic value (IG, GR, Gini, ...) is selected (best of them)
- [info](#)
 - for classification trees: usually, the most frequent class
 - for regression trees: usually, the mean of all target values in that leaf
 - * possible: median
- [stop_criterion](#)
 - 1. cautious splitting, post-pruning
 - 2. threshold
 - * classification: all instances have the same class, pure
 - * regression: variance
 - too few examples to continue splitting
 - * only allow tests that yield subtrees with at least two examples each
- impurity measures
 - good impurity measures are strictly concave <- **Question: do not understand!!**

2.6 Computational complexity

- Given a data set D consisting of N instances (x,y), where x has m components (attribtues), how do N and m influence the computational effort required to learn a tree?
- "Splitting one node" - for each node, we need to "find the best test"
 - for each possible test, evaluate its quality

- * partition the dataset according to this test
 - * compute quality based on this partition
- for the test with highest quality
 - * partition data according to that test
- Efficiently computing test quality for continuous attributes
 - first **sort** all tuples according to A (attribute), small to large
 - then gradually move the threshold, and simply update the tables
 - thus, testing **all thresholds** is only slightly more work than testing one!
- Computing the quality of one test
 - for nominal attributes: 1 scan, gradually building the table; compute IG once, at the end
 - for continuous attributes: 1 scan, gradually updating the table; compute IG for each intermediate table
 - hence, computing the best test for a node is linear in the number of instances in that node $O(n)$, with n the number of instances
 - Complexity:
 - * compute quality of one test is $O(n)$
 - * compute quality of all tests is $O(nm)$, with m the number of attributes
- Splitting multiple nodes
 - the total at one level of the tree is always n, so the overall amount of work depends on the layer of the trees
 - * if the splits are balanced, the height of tree is $O(\log(N))$, and the overall complexity of tree growth is $O(mN \log_2(N))$
 - * if the splits however are unbalanced, the height of tree is $O(N)$, and the complexity is $O(mN^2)$
- For impurity measures
 - "pure" = "zero variance/entropy"
- Why decision tree for Big Data?
 - QUICK! and FAST! efficient!
 - high accuracy
 - interpretable can be also a reason

2.7 Missing values

- when computing quality of a test: just ignore instances where this value is missing
- when splitting on an attribute:
 - guess its most likely value
 - partially assign the example to multiple branch, with its probability

2.8 Model trees

Model trees are **regression tree** in which each leaf does not contain a constant, but a linear model.

- RETIS (M5)
- Mauve

2.9 Multi-target trees

- Classification and regression trees are very popular for predicting one target variable
- But the principle of variance reduction is easily generalized to predicting vectors!
- This allow us to predict multiple variables at the same time, using one tree. (**vector**)
- Multi-label classification
 - How? For example for $Y = \{a,b,c,d,e\}$
 - Option 1: "binary relevance"
 - * learn (y/n) decision tree for each label
 - * TreeA predicts $a(y/n)$, TreeB predicts $b(y/n)$
 - Option 2: "label powersets"
 - * consider each set as separate label, for 5 original labels a,b,c,d,e, we get 32 combined labels: -,a,b,c,d,e,ab,ac, ..., abc, ...
 - * Learn a tree that predicts the combined label
 - Option 3: "Vector encoding"
 - * encode each set as 0/1 vector, e.g. $\{a,b,d\} = [1,1,0,1,0]$
 - * use a learner that can learn models that predict vectors
 - Option 3: "Vector encoding" is better! Requires almost no changes in the algorithm.

- * for example a leaf contains $\{[1,1,0,1,0],[1,1,0,0,0],[1,1,0,1,1]\}$ -> $[1,1,0,0.67,0.33]$
 - * predict all labels by threshold (e.g. 0.5): $[1,1,0,0.67,0.33]$ -> $\{a,b,d\}$
- Hierarchical multilabel classification (HMC)
 - for protein
 - similar just add a constraint "a label can only occur in an instance's label set if all its ancestors also occur"
 - * 250 functions = 250 trees is not fast and interpretable!
 - * 250 functions = 1 tree is fast and interpretable!
 - * learn 1 tree that predicts 250-dimensional class vector
 - Decision tree can be converted to rules set

2.10 Practical software

- Weka
- PythonL scikit-learn
- others: R, SAS, SPSS

3 Lecture 3: Learning sets of rules

3.1 Linear regression

Decision tree vs. linear regression

- "Linear model": $Y = a + b_1X_1 + b_2X_2 + \dots + b_kX_k$
 - usually fit such that sum of squared vertical deviations from line is minimal
 - what can we learn from such linear model?
 - * for predicting Y, given X_i
 - * for understanding how well Y can be predicted from X_i
 - * for understanding what the effect of each X_i is on Y
 - * for visualizing the connection between X_i and Y
 - * the linear correlation r tells us how well the points fit a line $-1 \leq r \leq 1$
 - * the coefficient of determination R^2 tells us to what extent Y is determined by the X_i , $0 \leq R^2 \leq 1$
 - careful with interpretation of coefficients
 - * coefficients are not scale-free
 - * "multicollinearity": correlations among X_i
- important assumptions
 - effect of each variable on target is constant (does not depend on other variables)
 - effect of different variables are cumulative (add up)
- complex terms
 - "overall" coefficient of X_2 is $(b_2 + b_{12}X_1)$ effect of X_2 on Y depends on X_1
- nominal variables
 - for nominal X_i with k values, introduce k-1 "0/1" variables, called "indicator" or "dummy" variables <- **Question: do not understand the dummy!!!**

3.2 Trees vs. linear regression vs. inductive bias

- Each learning approach has a "bias": implicit assumptions it makes.
- Removing all bias?
 - bias-free learning is impossible!

- no single best method for learning!
- for VS, the bias is "it assumes conjunctive concepts". -> without bias, no generalization.
- all of learning models have their own bias = implicit assumptions about what properties the true model has
- Choices to make
 - Modelling your problem as a prediction tasks:
 - * What is input, what is output (target attribute)?
 - * Regression, classification, probability prediction?
 - Choosing a learning approach
 - * Efficiency of learning/prediction phase
 - * Bias (which more fits the problem)
 - * interpretability of returned models (interpretation)

3.3 Rule learning

Learning sets of classification rules: rule sets:

1. "if...then..."
 2. "if...then...else..."
- Example: rule sets - define a leap years
 - If year is multiple of 400 then leap
 - else if year is a multiple of 100 then not leap
 - else if year is multiple of 4 then leap
 - else not leap
 - A decision tree can be turned into a set of rules!
 - By learning decision tree to learn rule sets
 - principle
 - 1. High accuracy: when it makes a prediction, it should be correct
 - 2. Reasonable coverage: it needs not make a prediction for each instance, but the more, the better
 - Could be top-down or bottom-up
 - Top-down:
 - * Start with maximally generally rule
 - * add literals one by one

- * gradually maximize accuracy without sacrificing coverage
- Bottom-up:
 - * Start with maximally specific rule
 - * remove literals one by one
 - * gradually maximize coverage without sacrificing accuracy

```

function LearnRuleSet(Target, Attrs, Examples, Threshold):
  LearnedRules :=  $\emptyset$ 
  Rule := LearnOneRule(Target, Attrs, Examples)
  while performance(Rule, Examples) > Threshold, do
    LearnedRules := LearnedRules  $\cup$  {Rule}
    Examples := Examples \ {examples classified correctly by Rule}
    Rule := LearnOneRule(Target, Attrs, Examples)
  sort LearnedRules according to performance
  return LearnedRules

```

Figure 2: Algorithm for "General algorithm for rule learning"

```

function LearnOneRule(Target, Attrs, Examples):
  NewRule := "IF true THEN pos"
  NewRuleNeg := Neg
  while NewRuleNeg not empty, do
    // add a new literal to the rule
    Candidates := generate candidate literals
    BestLit :=  $\text{argmax}_{L \in \text{Candidates}}$  performance(Specialise(NewRule, L))
    NewRule := Specialise(NewRule, BestLit)
    NewRuleNeg := {x  $\in$  Neg | x covered by NewRule}
  return NewRule

function Specialise(Rule, Lit):
  let Rule = "IF conditions THEN pos"
  return "IF conditions and Lit THEN pos"

```

Figure 3: Algorithm for "General algorithm for learning one rule"

- Top-down: start with an empty rule
- Heuristics for rule learners
 - High accuracy (most important), it is not robust to noise

- reasonably high coverage
- if-then-else rules vs. decision lists
 - if-then-else rules
 - * if year is a multiple of 400 then leap
 - * else if year is a multiple of 100 then not leap
 - * else if year is multiple of 4 then leap
 - * else not leap
 - decision lists
 - * if year is a multiple of 400 then leap
 - * if year is multiple of 4 but not of 100 then leap
 - if-then-else rule is more interpretable
 - decision list is more compact
 - unordered rules vs. ordered rules <- **Question: do not understand!!!**
- example-driven top-down rule induction
 - works like regular top-down approach, except:
 - * pick a not-yet-covered example
 - * consider as hypothesis space, all the rules that cover this example
 - * search within this hypothesis spaces (much smaller)
 - **pros:** more efficient
 - **cons:** less robust to noise
- other examples: RIPPER, Weka (software)

3.4 Association rules

Table 1: Differences with Classification and association rules

Classification rules	Association rules
One target class	Any combinatino of items can be the target
Good rules have near 100% accuracy ("confidence" here)	Rules need not have near 100% confidence to be interesting
Find a minimal set of rules (just enough to classify)	Find a maximal set of rules (all rules that hold)

- Overview
 - Similar to classification rules, but for descriptive learning instead of predicitive learning
 - is to look for the patterns in data

- classification rules are a small subset of association rules?
- General format: if a_1, a_2, \dots, a_n then $a_{n+1}, a_{n+2}, \dots, a_{n+m}$
- Rule "If <this> then <that>" is characterized by
 - Support: % of all clients that buy <this>
 - Confidence: % of buyers of <this> that also buy <that>
- Running a classification rule learner clearly will not work well for association rule, since classification rule is only a small subset of association rule
 - Solution: APRIORI algorithm

4 Lecture 4: Instance-based learning, Clustering

4.1 Instance-based learning

Basic key idea: just store all training examples

- When seeing a new instance:
 - Find the most similar cases in the database
 - make a prediction based on those instance
- Architypical method: k-nearest-neighbors (k-NN)
 - use most frequent class/ mean target among the k nearest neighbors as your prediction
 - k is chosen by the user (hyperparameter)
- Similarity
 - how close to others, often Euclidean distance for numerical inputs
- Voronoi diagrams
 - indicate area where prediction is influenced by same set of examples
 - for 1-NN: cell borders are right in the middle between any two data points
 - This is called a Voronoi diagram - kNN
- Decision surface
 - Decision surface separate regions with different predictions
- Voronoi diagram for $k > 1$
 - To construct diagram for k-NN with $k > 1$
 - * Start from diagram for k-1
 - * for each cell:
 - temporarily forget about k-1 nearest neighbors
 - split cell according to k'th nearest neighbors
 - * Merge adjacent cells with same k nearest neighbors
- pros vs. cons
 - **pros**
 - * "Learning" is very fast - just storing the data
 - * all detials of the data are kept
 - **cons**

- * can be slow at prediction time
 - * difficulties in high-dimensional spaces
 - * relies on having a good similarity measure (Euclidean distance - numerical)
 - * not robust to noise
- for k-NN, large K -> more robust to overfitting? but it is more robust to noise
- difficulties with high-dimensional space - curse of dimensionality - data are distributed very sparsely in high-D
- Improvement: Different scales
 - When dimensions have very different scales, Euclidean distance may not work well
 - * Solution: normalization - normalize all dimensions to comparable scale
 - * Give irrelevant dimensions a smaller weight in the Euclidean distance, so they have less influence
 - * using a closed formula: $w_i = 1 - \frac{1}{n} \sum_{k=1}^c \sum_{j=1}^{n_k} |\bar{x}_{ki} - x_{kji}|$
 - * with c =#clusters, n_k =#elements in cluster k, \bar{x}_{ki} =mean x_i in cluster k, $x_{kji}=x_i$ for jth element in cluster k
- Improvement: distance-weights k-NN
 - Why 3-NN, but why not 4-NN
 - * Solution: no cut-off at k, but have weights gradually decrease with distance
 - * careful: influence must decrease fast with distance, otherwise faraway cases will dominate the voting (otherwise noise)
- Improvement: locally weighted regression
 - for better fitting
 - * Solution: fit a simple local model
 - * a linear regression can be okay
- Prototypes
 - A prototype is a representative for a group of instances
 - can be an average for elements in a cluster
 - note: need to define similarity between instance & prototype
- Efficient prediction
 - Solution: indexing the data helps

- Lazy learners vs. eager learners
 - k-NN is called lazy learner
 - * do not build a model during training, merely store data
 - * start doing the hard work when asked a question
 - eager learners
 - * generalize before knowing the question
 - * try to build a global model that will work under all circumstances
 - lazy learners can use simpler local models - sometime very accurate
 - * compare: locally weighted linear regression (k-NN) vs. global linear regression
 - * compare: learn one rule that covers the query instance vs. learn a global rule set

4.2 Clustering

- Overview
 - Find structure/pattern in the data, in the form of groups of instances that are highly similar to each other
 - It is unsupervised machine learning (no labeled samples)
- Find groups (clusters) of instances so that
 - instances in the same group are similar
 - instances in the difference group are different
- Definition of clustering
 - flat clustering
 - * returns a partition of the data
 - hierarchical clustering
 - * returns a hierarchy of clusters
- Other definition of clustering
 - extensional clustering
 - * clusters are defined without any description language
 - conceptual clustering
 - * clusters are defined using a conceptual description language
- FFlat, extensional clustering

Input : dataset $D \subset X$, number of clusters k Output : partition of D into k clusters Algorithm: Choose k random <i>seeds</i> (points in X) Repeat until no changes: Assign each instance to the cluster of its closest seed Redefine seeds as cluster means Return the k clusters
--

Figure 4: Algorithm for "Algorithm for flat, extensional clustering: K-means"

- task: given a set of unlabeled data, find groups of highly similar instances
- some constraints may additionally be given
- procedure: reassign points, recompute seeds, reassign points ..., when no points are re-assigned to another cluster, stop.
- Hierarchical, extensional clustering
 - top-down ("divisive") methods
 - * start with 1 cluster (whole data set)
 - * divide it into subsets
 - * subdivide subsets further
 - bottom-up ("agglomerative") methods
 - * start with singleton clusters
 - * join closest clusters together
 - * repeat until 1 cluster
- Bottom-up methods
 - How does it generalize to distance between clusters? - based on examples
 - * Single linkage: distance between clusters = distance between each closest point in their cluster
 - * Complete linkage: distance between clusters = distance between each furthest point in their clusters
 - * Average linkage: distance between clusters = distance between average of clusters
- Conceptual clustering
 - Cluster + find a conceptual description of each cluster (in some given description language L)

- Clearly, the clusters are not defined using distance alone! Context is important!
- Examples: Cobweb
 - Predictability: given cluster, how well can you predict attribute values
 - Predictiveness: given attribute values, how well can you predict cluster
 - Cobweb is to maximize a combination of both
- Decision tree learning, viewed as clustering
 - A decision tree defines a conceptual, hierarchical clustering of the data
 - * each node = subset of the data
 - * conceptual description of these data = conjunction of test outcomes from root to node
- Clustering trees
 - for regression (original), heuristic minimize average variance of Y within subsets
 - * $Var(S) = \sum_{(x,y) \in S} \frac{(y-\bar{y})^2}{|S|-1}$
 - * with $\bar{y} = \sum_{(x,y) \in S} \frac{y}{|S|}$
 - for clustering (unsupervised clustering tree), given some distance metric d that indicates dissimilarity, we can minimize the average variance of X within subsets
 - * $Var(S) = \sum_{x \in S} \frac{d(x, \bar{x})^2}{|S|-1}$
 - * with $\bar{x} = \sum_{x \in S} \frac{x}{|S|}$
- Predictive clustering
 - overview: it looks like from value \rightarrow cluster and from cluster \rightarrow value
 - Predictive clustering builds a model consisting of
 - * a set of clusters
 - * a function c assigning instances to clusters
 - * a function p assigning a target value to the instance, given the cluster
 - the overall predictive function is then $f(x) = p(c(x), x)$
 - the accuracy of f will depend on that of c and p.
 - * an accurate c requires good predictiveness (from value to cluster)
 - * an accurate p requires good Predictability (from cluster to value)




Figure 5: Algorithm for "Scheme of Predictive Cluster"

- Similarity measures
 - Similarity is often represented using a distance metric
 - * small distance = high similarity, e.g., euclidean distance, manhattan distance, hamming distance, ...
 - * $d_{Eucl}(x, x') = \sqrt{\sum_i (x_i - x'_i)^2}$
 - * $d_{Manh}(x, x') = \sum_i |x_i - x'_i|$
 - Not all similarity measures can be expressed in that way!
 - * Distance metric fulfills symmetry and triangle inequality
 - * Some similarity measures cannot be mapped to such a distance measures
- Learning the similarity measure
 - Clustering relies strongly on defining an appropriate similarity measure
 - However, this may be very difficult sometime
 - * Solution: semi-supervised clustering
 - * Computer learns the similarity from examples of pairs of instances that should be in the same/ in different clusters, "must-link" and "cannot-link" constraints
 - * then start the unlabeled data, (semi-supervised clustering: a few labeled data, many unlabeled data)
- Learning preferred clusterings
 - Different clustering algorithms have different biases
 - * k-means: spherical clusters
 - * density-based: can return concave clusters
 - * some methods can even return disconnected cluster
- example of interactive clustering: the COBRAS method

- using an intermediate level of super-instances
 - * clustering = set of clusters
 - * cluster = set of super-instances
 - * super-instance = set of clusters
- Evaluating clusterings
 - How can we assess whether a clustering is good or bad?
 - explicit objective: minimize intra-cluster variance
 - internal criteria: inherent to the clustering
 - external criteria: compare to a reference clustering
 - * rand (random) index
 - * adjusted rand (random) index (ARI)

5 Lecture 5: Evaluating hypotheses

5.1 Models and learning algorithms

- Different levels of evaluation
 - evaluation of learned models
 - * given the model, how well can we expect it to perform
 - evaluation of learning algorithm
 - * given an algorithm, how well can we expect the models it learns to perform
- Evaluation of classifiers
 - probability of making a correct prediction: accuracy
 - time needed to make the prediction
 - cost incurred by wrong predictions

5.2 Statistics

- accuracy
 - accuracy = probability of correct prediction on a randomly drawn instances
 - error = 1- accuracy
 - estimating accuracy = estimating a probability over a population
 - estimating probability from sample proportion: well-studied problem in inferential Statistics
 - for probability π , sample proportion p has $\mathbb{E}(p) = \pi$ and $Var(p) = \frac{\pi(1-\pi)}{n}$
 - Hence:
 - * point estimate: $\hat{\pi} = p$
 - * $1-\alpha$ confidence interval: $[p - z_{\alpha/2}\sqrt{\frac{p(1-p)}{n}}, p + z_{\alpha/2}\sqrt{\frac{p(1-p)}{n}}]$ with $z_{\alpha/2}$ such that $P(Z > z_{\alpha/2}) = \alpha/2$ if $Z \sim N(0, 1)$

Table 2: Alpha for Confidence interval

α	0.1	0.05	0.01
$z_{\alpha/2}$	1.64	1.96	2.56

- train set and cross-validation
- option 1: train/test split

- use 2/3 of available data for training set T, set aside 1/3 for test set S
- while if T is smaller \rightarrow less accurate f, smaller S \rightarrow less accurate estimate of accuracy of f
- option 2: cross-validation
 - learn f from the full data set S
 - however to get an estimate of $\text{acc}(f)$, we partition the set into n subsets S_i
 - learn f_i , $i=1\dots n$ from $\frac{S}{S_i}$, compute $\text{ACC}(f, S_i)$
 - estimate $\text{acc}(f)$ as average of all $\text{acc}(f, S_i)$
 - name
 - * for a specific n, this is called n-fold cross-validation
 - * when $n = |S|$, this is called leave-one-out cross-validation
 - Is cross-validation entirely unbiased - No
 - More option: nested/internal cross-validation
- comparing two models
 - given two models, which one has higher accuracy
 - two cases
 - * compare 2 models on different test sets
 - * compare 2 models on same test sets
 - different sets
 - * Checked by confidence interval: $(a_1 - a_2) \pm z_{\alpha/2} \sqrt{\frac{a_1(1-a_1)}{n_1} + \frac{a_2(1-a_2)}{n_2}}$
 - * if CI for $a_1 - a_2$ is entirely to the right of 0, meaning that $a_1 - a_2 \geq 0$, meaning that $\text{acc}(f_1) \geq \text{acc}(f_2)$
 - * if CI for $a_1 - a_2$ is entirely to the left of 0, meaning that $a_1 - a_2 \leq 0$, meaning that $\text{acc}(f_1) \leq \text{acc}(f_2)$
 - * if CI contains 0, it means that there is no difference between f_1 and f_2
 - same sets
 - * Compare models on the same data set is more informative
 - * uses more detailed information from test
 - * use McNemar's test in this case
 - * key idea: how often was f_1 right and f_2 wrong on the same example vs. how often was f_1 wrong and f_2 right on the same example
 - if $B \approx C \approx (B + C)/2$, then accept $\text{acc}(f_1) = \text{acc}(f_2)$

- otherwise, if B deviates too much from $(B+C)/2$, then reject $acc(f_1) = acc(f_2)$
- * McNemar's test is more informative - but can only use it when individual predictions of both models on the same test set are available

Table 3: Table for McNemar's test

	f_1 right	f_1 wrong
f_2 right	A	B
f_2 wrong	C	D

- Evaluation based on accuracy is not always appropriate
- cons
 - no obvious reference point
 - * sample imbalance
 - unstable when class distribution may change
 - * because classifier has different TP and TN
 - *assumes symmetric misclassification costs*
- alternatives
 - correlation
 - ROC analysis
- correlation
 - correlation tends to be more informative for unbalanced classifiers
 - $\phi = \frac{AD-BD}{\sqrt{T_+ \cdot T_{-} \cdot T_{pos} \cdot T_{neg}}}$ with 1: perfect prediction, 0: no correlation (totally random), -1: predicting the exact opposite (1 and -1 both indicate strong relation, more focus on absolute number)

Table 4: Real class vs. predicted class

	+	-	Total
pos	A	B	T_{pos}
neg	C	D	T_{neg}
Total	T_+	T_-	T

- expected misclassification cost

- sometimes one type of mistake is worse than the other, accuracy ignores this (assumes all errors are equally bad -> which are not correct usually)
- Solution
 - * accuracy: probability that some instance is classified correctly
 - * TP: "true positive rate": probability of a positive to be classified correctly
 - * TN: "true negative rate": probability of a negative to be classified correctly
 - * FP: "false positive rate": $FP = 1 - TN$.
 - * FN: "false negative rate": $FN = 1 - TP$.
- consider "misclassification costs"
 - * C_{FP} : cost of a false positive (cost of classifying a - as pos)
 - * C_{FN} : cost of a false negative (cost of classifying a + as neg)
- expected cost of a single prediction
 - * $C = C_{FP}P(pos|-)P(-) + C_{FN}P(neg|+)P(+)$
 - * estimated by $C = C_{FP} \cdot FP \cdot T_-/T + C_{FN} \cdot FN \cdot T_+/T$
- ROC (Receiver operating characteristic) analysis
 - Allow us to see
 - * how well a classifier will perform given certain misclassification costs and class distribution
 - * in which environments one classifier is better than another
 - Overview: a ROC diagram plots (estimated) TP vs. FP for a classifier
 - * $TP = \frac{A}{T_+}$
 - * $FP = \frac{B}{T_-}$
 - * why T_+ and T_- instead of T_{pos} and T_{neg} ?
 - 1 classifier = 1 point in the ROC diagram
 - * higher is better (better TP), more to the left is better (lower FP)
 - * TP = 1: no positives forgotten
 - * FP = 0: no negatives can pretend positives
- Rank classifier
 - a rank classifier returns a numerical score, rather than pos/neg
 - * expresses uncertainty: higher score = more certain of pos
 - * numerical score can be turned into binary prediction by choosing a threshold
 - * changing the threshold gives different points in the diagram, forming ROC curve

- Iso-cost line
 - given cost c_{FP} and c_{FN} , the expected cost of a classifier with TP and FP is $c = c_{FP} \cdot FP \cdot \frac{T_-}{T} + c_{FN} \cdot (1 - TP) \cdot \frac{T_+}{T}$
 - hence, for a **constant cost c**, the points with the cost are on the following line - called an iso-cost line:
 - * $TP = \frac{c_{FP}T_-}{c_{FN}T_+} + 1 - \frac{cT}{c_{FN}T_+}$
- Convex hull
 - set of classifiers = set of points in ROC diagram
 - classifiers that are optimal under some condition are on the convex hull of this set (if can be compared, highest is the optimal, other is not on the convex hull)
- precision-recall diagram
 - like ROC, defined for binary (pos/neg) predictions
 - * Precision: what proportion of positive predictions is really positive? using $P(+|pos)$
 - * Recall: what proportion of positives is predicted positive? using $P(pos|+)$
 - like with ROC
 - * classifier -> point
 - * rank classifier -> curve

5.3 Comparison of evaluation

- Evaluation
 - evaluating classifier
 - * discussed above
 - evaluating regression models
 - * SSE, MSE, RE, correlation
 - evaluating clustering
 - * explicit objective, internal criteria, external criteria
 - evaluating learners
 - * two questions
 - Q1: given models f_1 and f_2 , which one has better predictive accuracy?
 - Q2: given learners L_1 and L_2 and data set S, which learner can be expected to build best model from (data set like) S?
 - * Q2 is more difficult to answer than Q1, additional level of variation

- * Q2 is only meaningful when restricted to a particular problem domain, and data set of a particular size
- * statisticians distinguish
 - Conditional accuracy: mean accuracy of models learned by L on a given dataset S
 - Unconditional accuracy: mean accuracy of models learned by L on datasets of size n drawn randomly from the population (from population)
- * comparing learners
 - making some rules, description language

6 Lecture 6: Numerical approaches (ANN, SVM), Computational learning theory

6.1 Learning theory

aaa

7 Lecture 7: Probabilistic approaches, Ensembles

8 Lecture 8: Reinforcement learning

9 Lecture 9-10: Inductive logic programming